COMPSCI 677 Spring 2022

Lab 3: Caching, Replication and Fault Tolerance

Team Members: Maoqin Zhu, Yixiang Zhang

Deployment on AWS

This could be a simple tutorial about how to deploy our online application on AWS Cloud. We are providing you enough details about the deployment configuration in this part.

Step 1- Install AWS CLI

Just follow the steps described in the following link:
https://docs.aws.amazon.com/cli/latest/userguide/getting-started-install.html

Step 2- Obtain AWS Crendentials.

For each UMass student, we can get our own credentials through AWS Academy. First we direct to this page, and just click “Start Lab” button as follows. It may take several minutes to start the AWS environment.

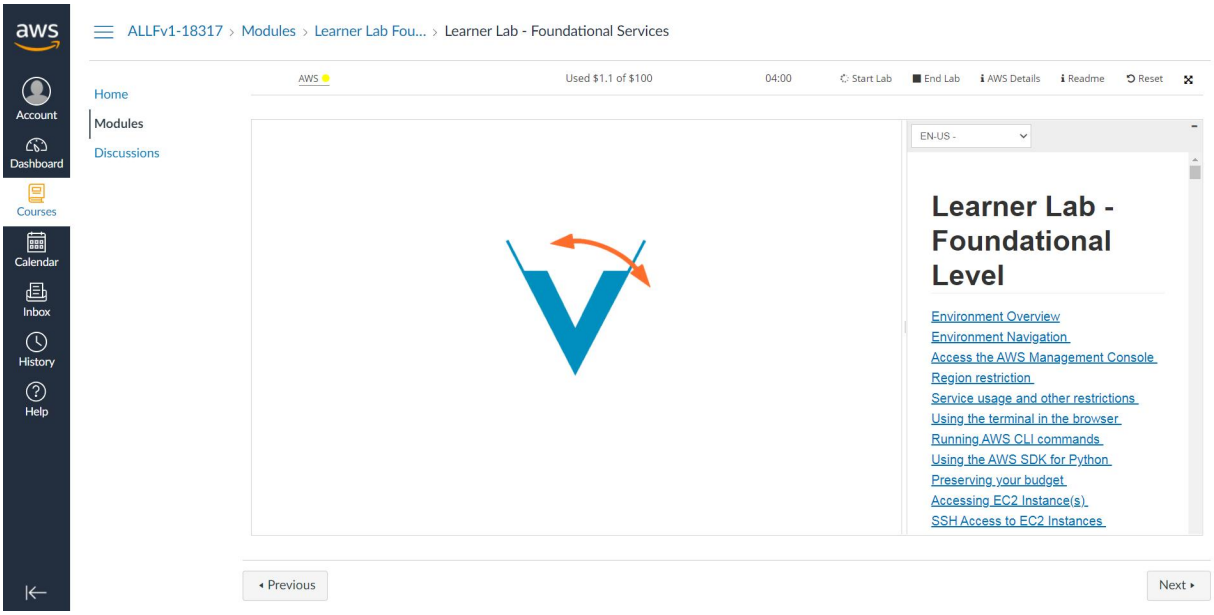

Click the "AWS Details" button, and then click the "Show" button following AWS CLI. The code that appeared in the text box below contains the credentials that the AWS CLI uses to interact with AWS.

What you should do on your Linux/macOS machine is that:

1) Copy the code and save it to **$HOME/.aws/ credentials**.

2) Download PEM key and save it to your working directory.

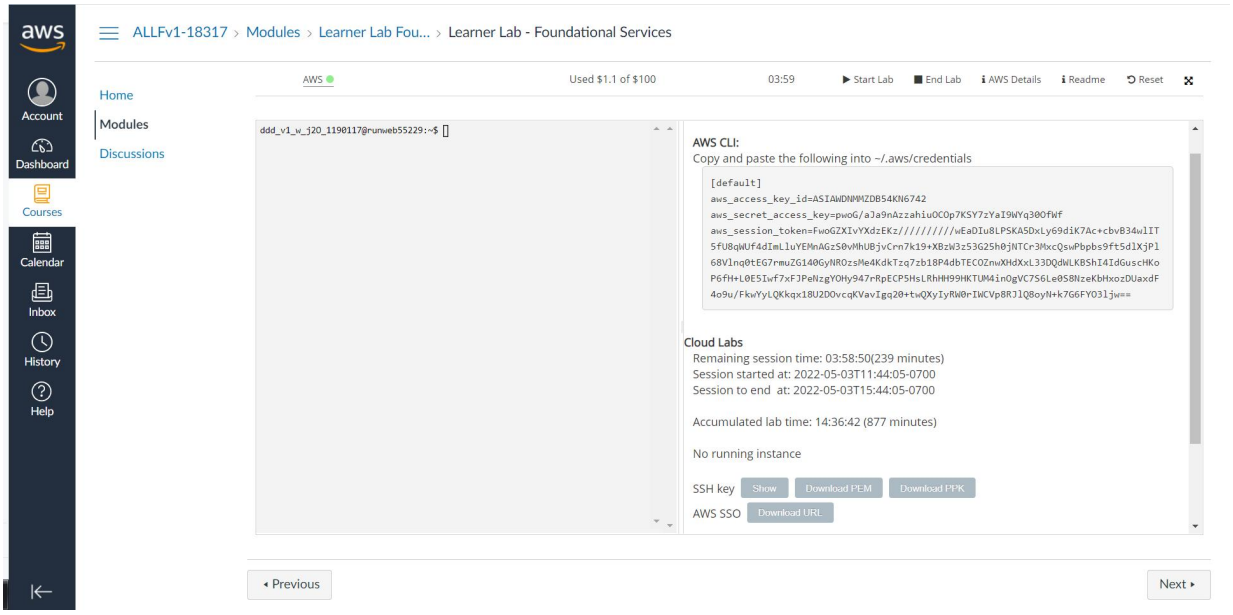

Step 3- Configure Your AWS Settings

Simply run **$aws configure** on your machine as follows. By creating a credentials file in previous step, we can just press enter to skip certain processes.

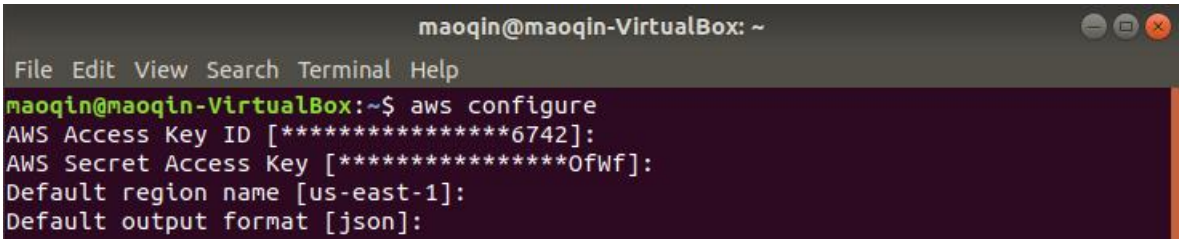

Step 4- Start Our EC2 Instance

When creating the EC2 instance, there are some options for you. Our selection in this lab are listed as below:

Instance Type: **m5a.large**

AMI ID: **ami-0d73480446600f555**

Hence we type the following command on our local machine.

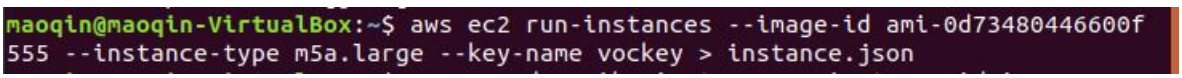

Step 5- **Checkout the Public IP**

First, type the following command on AWS Academy console:

$ aws ec2 describe-key-pairs

Then we will have an “instance.json” file on our local machine.

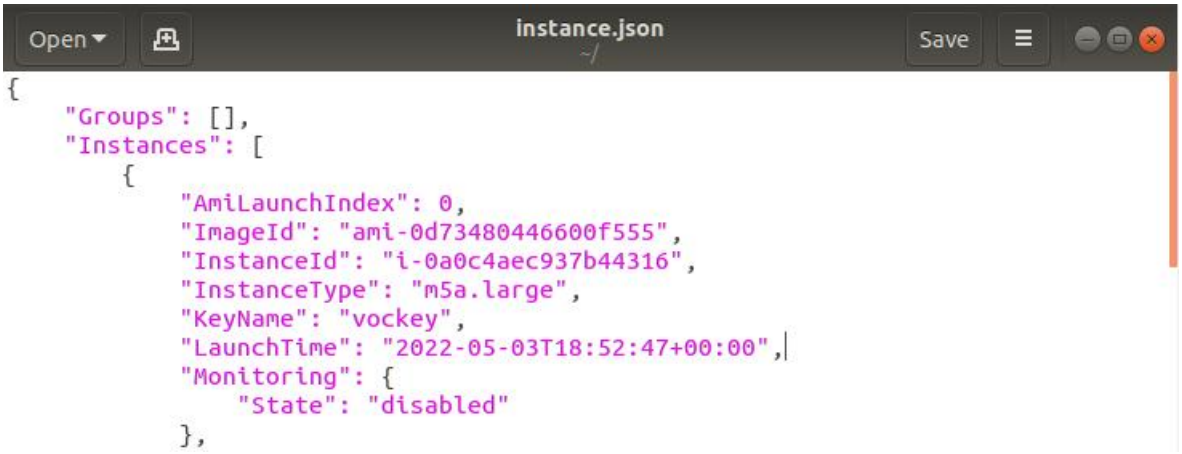

Second, according to our instance id, we can checkout the public IP address of our EC2 instance as follows:

$ aws ec2 describe-instances --instance-id <our-instance-id>

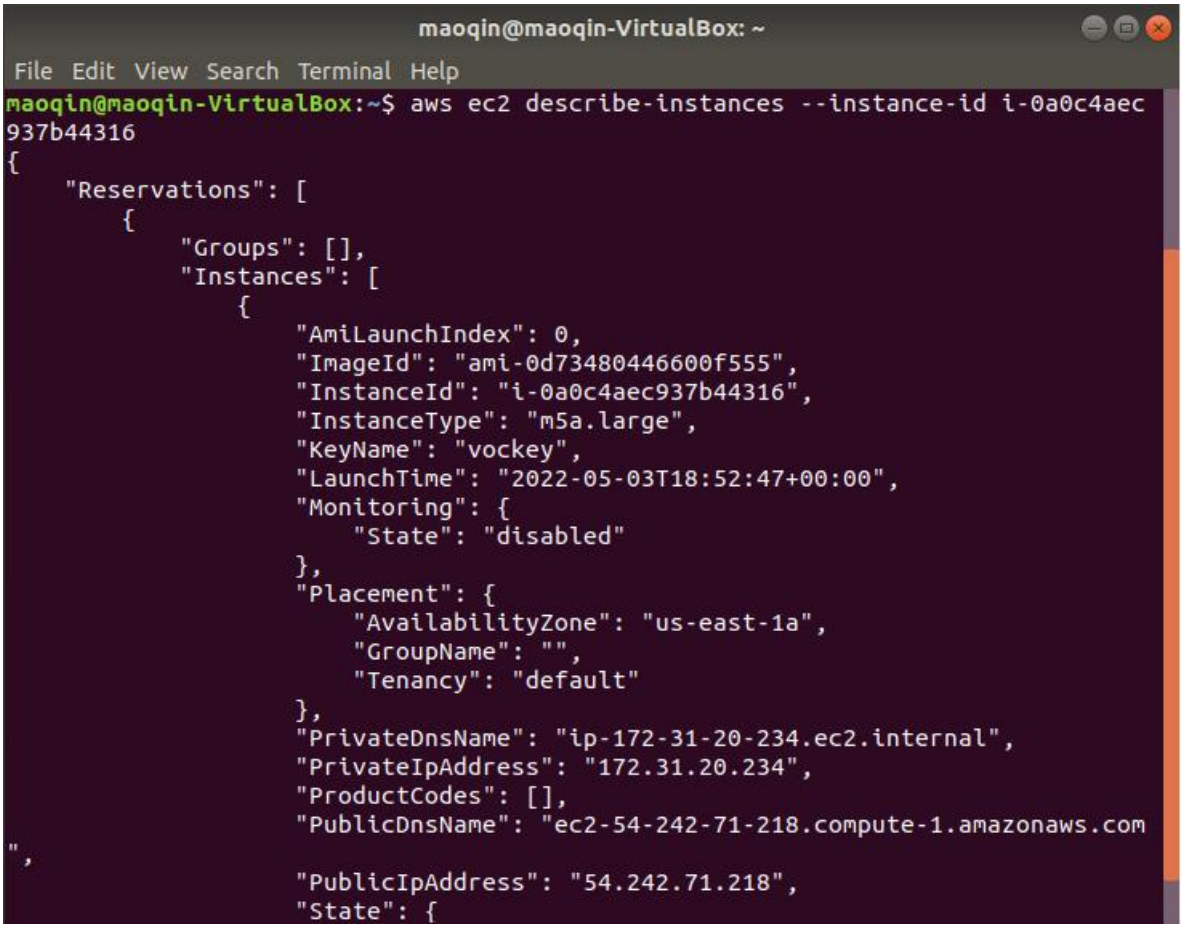

Step 6- Open Specific Ports on Our EC2 Instance

First of all, there are some ports to be opened in this lab:

SSH: **22**
Front-end Service: **6060**
Catalog Service: **10086**
Order Service: **10010-10012**

Type the following command on your local machine to open the ports:

$ aws ec2 authorize-security-group-ingress --group-name default --protocol tcp --port <number> --cidr 0.0.0.0/0

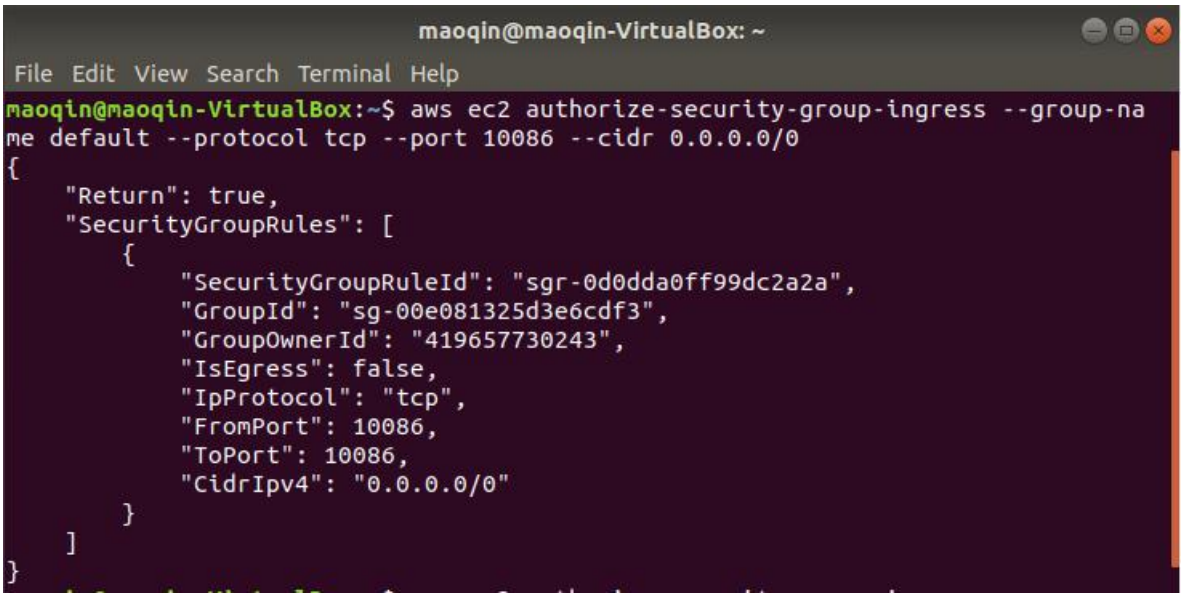

Step 7- Upload Our Source Code on EC2 Instance

In general, you can type command on your terminal to access our EC2 instance via SSH. Here we will use the PEM file we download previously for getting the permission. And then upload your source code on EC2 instance by Git.

Here we adopt the tools **XShell** and **Xftp7** to help us access our EC2 instance and upload source code via SSH. We are using the PEM file too.

Xftp7 View:

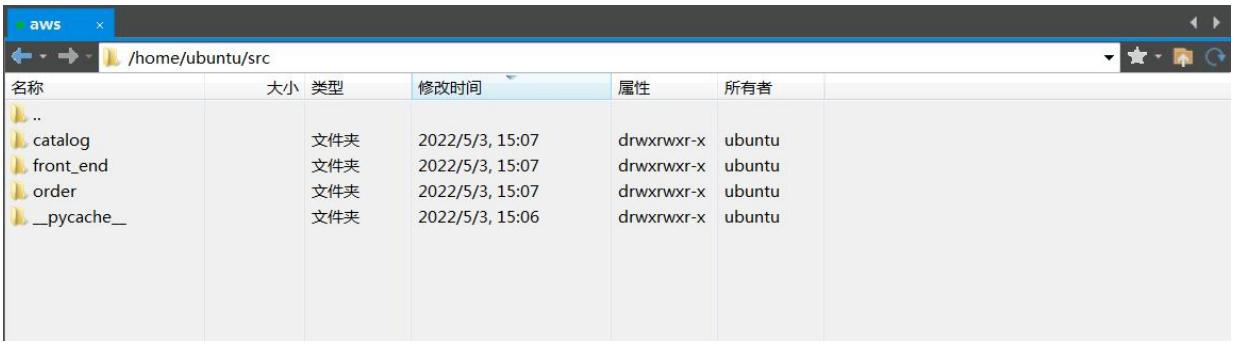

XShell View:

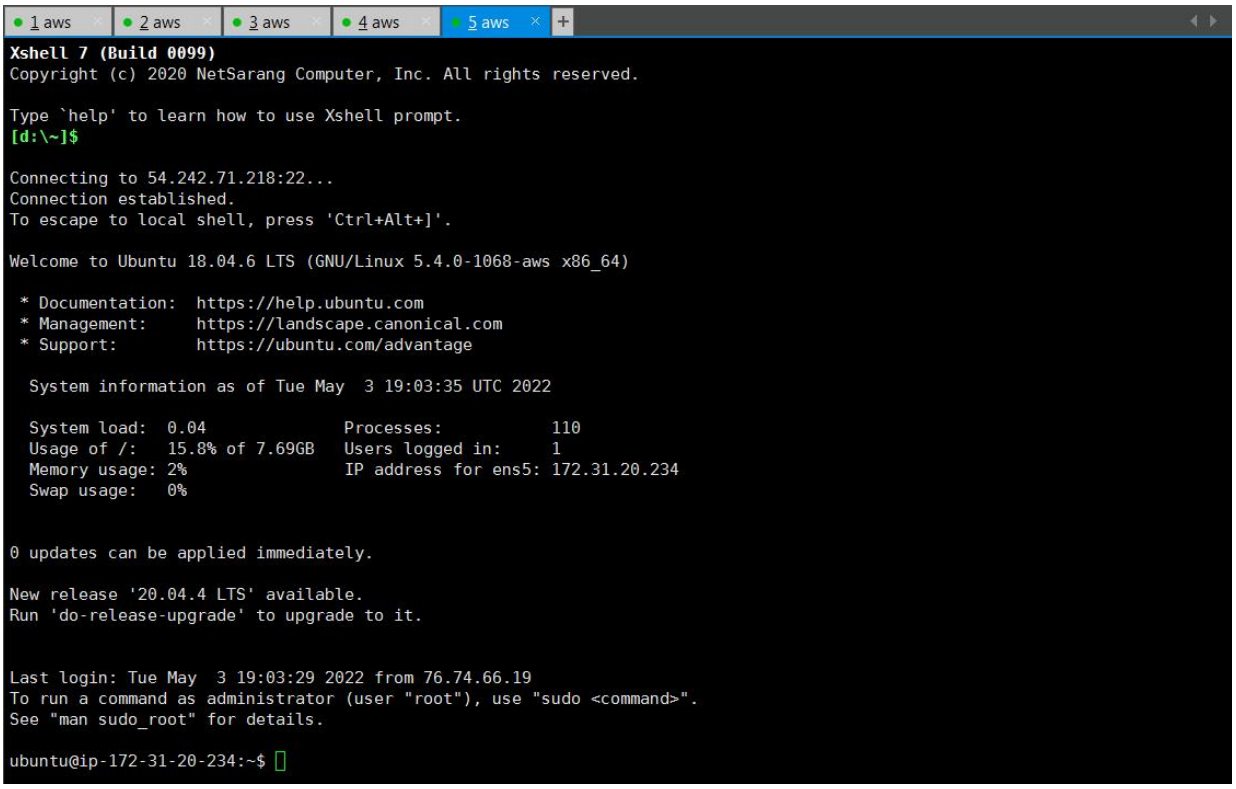

Step 8- Start Our Application

Notice that since we use several modules in our design, we recommend you run following commands on our EC2 instance at first:

Update System: **$ sudo apt-get update**
Install pip3: **$ sudo apt-get -y install python3-pip**
Install Flask: **$ pip3 install flask**

Now in order to start our application, simply open 5 terminals on our AWS EC2 instance, and run micro services in each terminal in specific order. After cd into different directories, type the following commands.

Terminal 1: **$ python3 catalog_server.py**
Terminal 2: **$ ID=1 PORT=10010 python3 order_server.py**
Terminal 3: **$ ID=2 PORT=10011 python3 order_server.py**
Terminal 4: **$ ID=3 PORT=10012 python3 order_server.py**
Terminal 5: **$ python3 front_end.py**

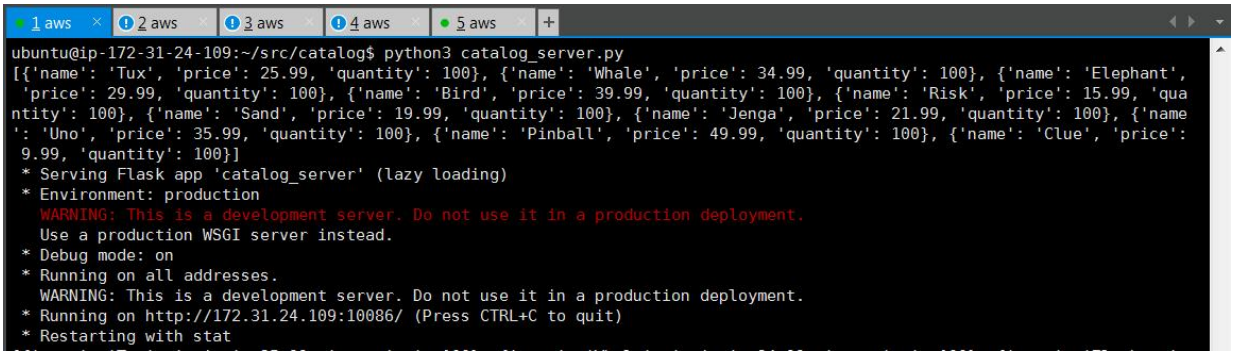

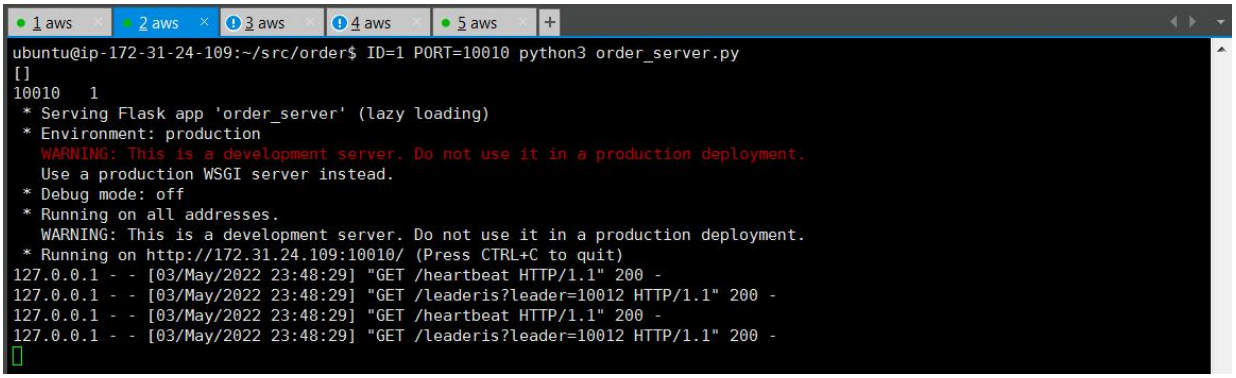
1 aws
2 aws
3 aws
4 aws
5 aws
ubuntu@ip-172-31-24-109:~/src/order$ ID=1 PORT=10010 python3 order_server.py
[]
10010 1
* Serving Flask app 'order_server' (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: off
* Running on all addresses.
WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://172.31.24.109:10010/ (Press CTRL+C to quit)
127.0.0.1 - - [03/May/2022 23:48:29] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [03/May/2022 23:48:29] "GET /leaderis?leader=10012 HTTP/1.1" 200 -
127.0.0.1 - - [03/May/2022 23:48:29] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [03/May/2022 23:48:29] "GET /leaderis?leader=10012 HTTP/1.1" 200 -

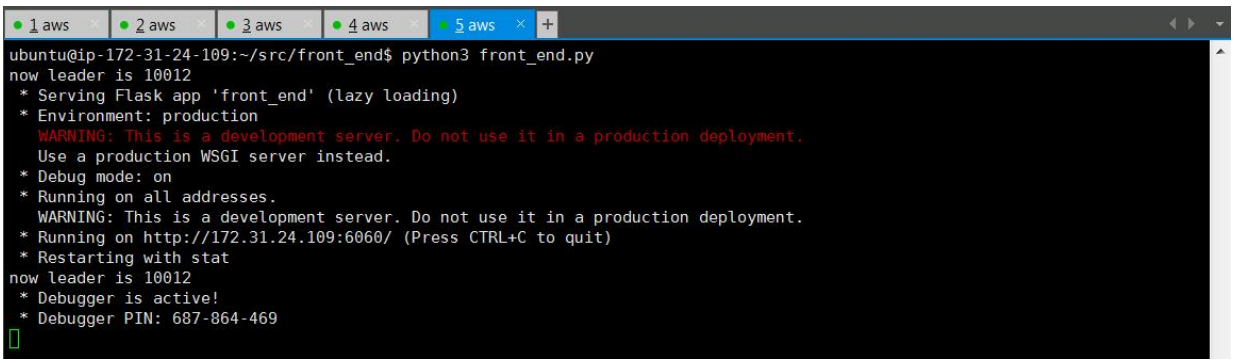
1 aws
2 aws
3 aws
4 aws
5 aws
ubuntu@ip-172-31-24-109:~/src/front_end$ python3 front_end.py
now leader is 10012
* Serving Flask app 'front_end' (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: on
* Running on all addresses.
WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://172.31.24.109:6060/ (Press CTRL+C to quit)
* Restarting with stat
now leader is 10012
* Debugger is active!
* Debugger PIN: 687-864-469

Testing & Evaluation

In terms of screenshots of our functional tests & load tests, please check out the **"output" file** for details. In this document, we mainly focus on analyzing the results we have.

1. Load Test Results

Our testing codes can automatically sends **1000** Query, Buy or queryOrder requests. Python unittest can help measure the total latency seen by clients in this case. Hence, in terms of average latency for each request, we should divide the total time by **1000**. For different type of requests, we repeatedly run 5 clients at the same time, and measure the total latency seen by each client.

Table1: Load test results

Request Type	*Average Latency / s*
Query	*0.07258*
Buy	*0.08021*
queryOrder	*0.06128*

2. Caching Test Results

In order to estimate how much benefits does caching. We are firstly measuring the latency seen by each client for different type requests with caching turned on. Change the probability p of a follow up buy request from 0 to 80%, with an increment of 20%, and record the result for each p setting. And then do the same experiments but with caching turned off.

For each experiment with different p, we are testing **multiple times**, and record the **average latency**. The final results for different p with caching turned on are shown as follows:

Table2: Test result with caching turned on

Probability	*Average Query Latency / s*	*Average Buy Latency / s*
0	*0.04702*	*n/a*
0.2	*0.04563*	*0.05703*
0.4	*0.04303*	*0.05764*
0.6	*0.04193*	*0.05713*
0.8	*0.04589*	*0.05772*

Then we make a plot showing the values of p on the X-axis and response time/latency on the Y-axis.

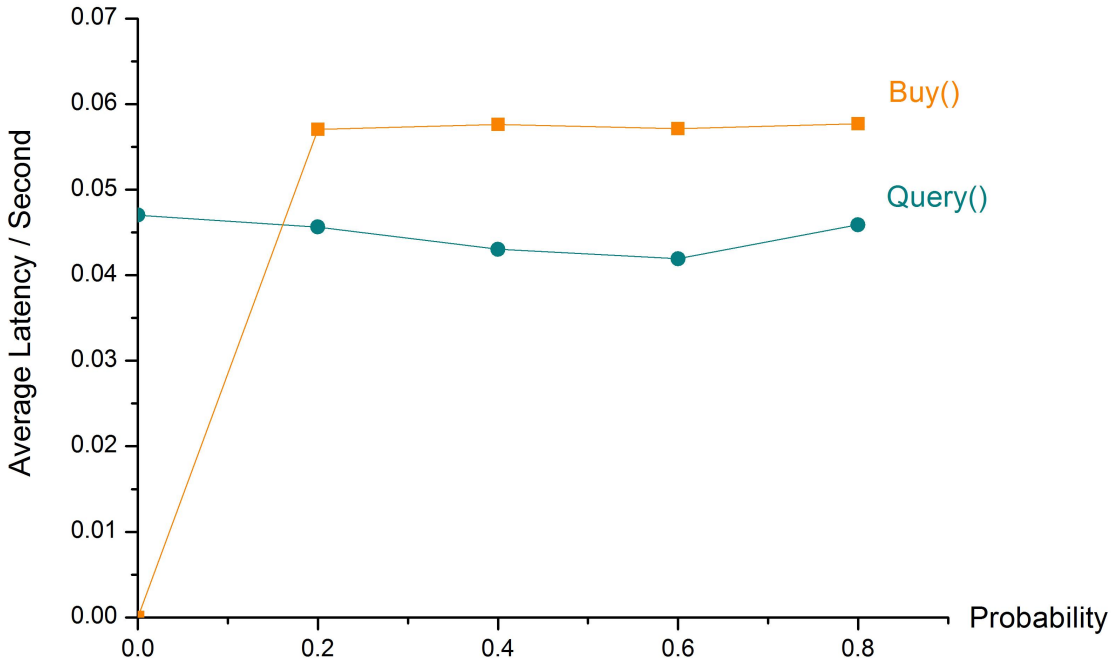

Figure1: Test result with caching turned on

And then just do the same experiments but with caching turned off. The final results for different p with caching turned off are shown as follows:

Table3: Test result with caching turned off

Probability	*Average Query Latency / s*	*Average Buy Latency / s*
0	*0.04754*	*n/a*
0.2	*0.04767*	*0.05917*
0.4	*0.04845*	*0.06005*
0.6	*0.04775*	*0.05976*
0.8	*0.04927*	*0.06240*

Then we make a plot showing the values of p on the X-axis and response time/latency on the Y-axis.

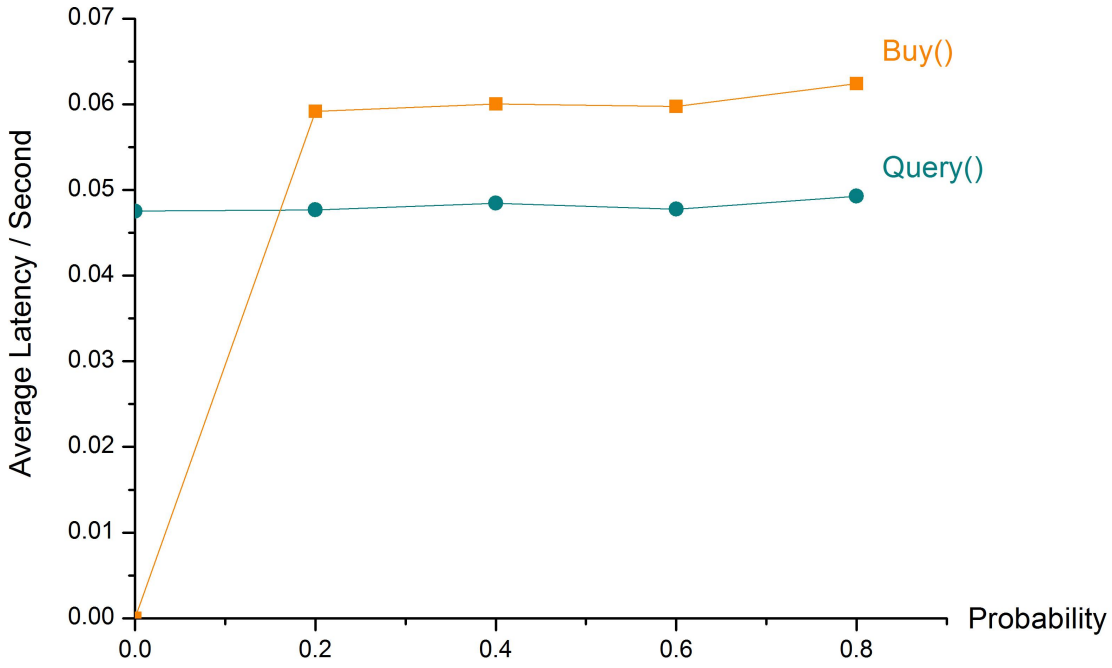

Figure2: Test result with caching turned off

Question 1

Estimate how much benefits does caching provide by comparing the results?

Solution:

As **table 2**, **figure 1**, **table 3** and **figure 2** shown above, the difference of average latency between different caching state (on/off) can be represented as follows:

Probability	*Query Latency Difference / s*	*Buy Latency Difference / s*
0	*0.00052*	*n/a*
0.2	*0.00204*	*0.00214*
0.4	*0.00542*	*0.00241*
0.6	*0.00582*	*0.00263*
0.8	*0.00338*	*0.00468*

Hence, we could say both the average query latency and the average buy latency with caching on is actually smaller than that with caching off.

Furthermore, the average benefits of caching for different types of requests are:

$$average\ benefits\ for\ query\ =\ 0.00344\ seconds\ better$$

$$average\ benefits\ for\ buy\ =\ 0.00297\ seconds\ better$$

3. Fault Tolerance Test Results

Finally, we are simulating crash failures by killing a random order service replica while the clients is running, and then bring it back online after some time.

Specifically, our **test case** follows the steps as described below:

Client Terminal: we are sending **1000** buy requests using the code in load test.

$ FRONT=<IP address> python3 -m unittest -v test_load.TestLoadPerformance.test_load_buy

Crash the follower with id = 1: terminate the node with **port = 10010** & **id=1**

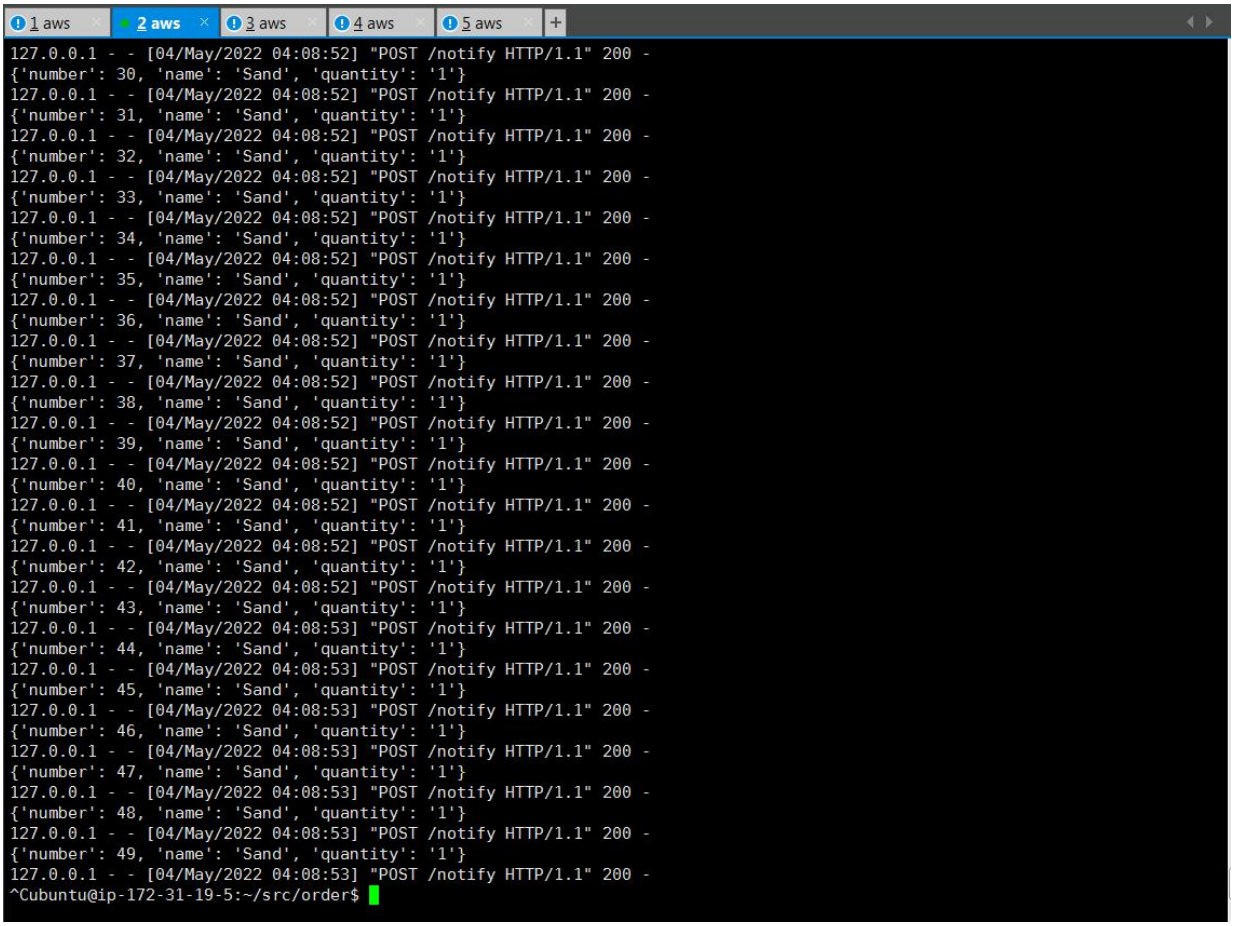

Restart the follower with id = 1: restart the node with **port = 10010** & **id=1**

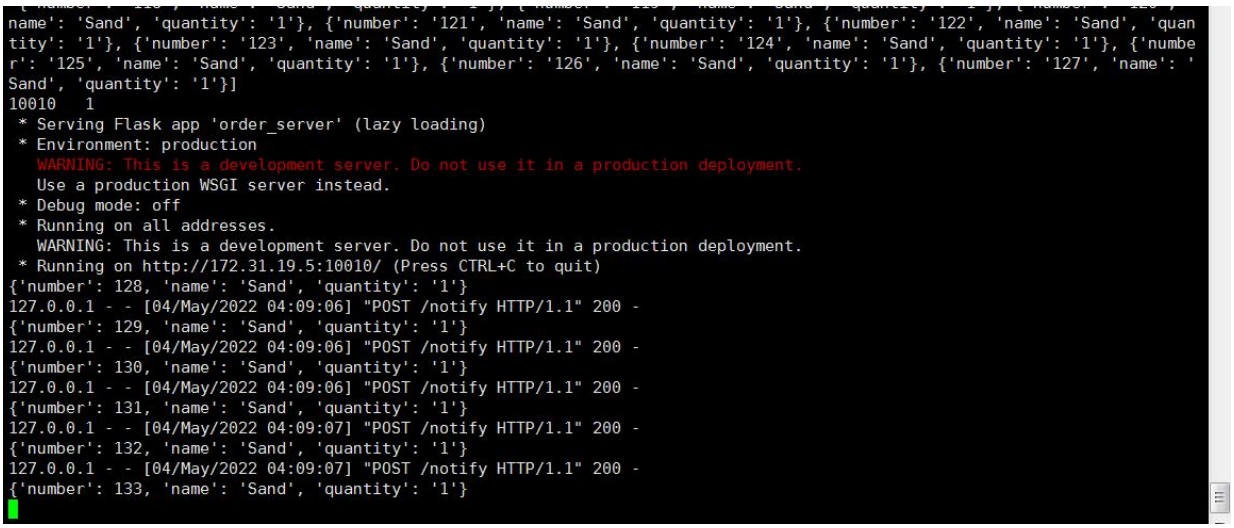

Crash the leader with id = 3: terminate the node with **port = 10012** & **id=3**

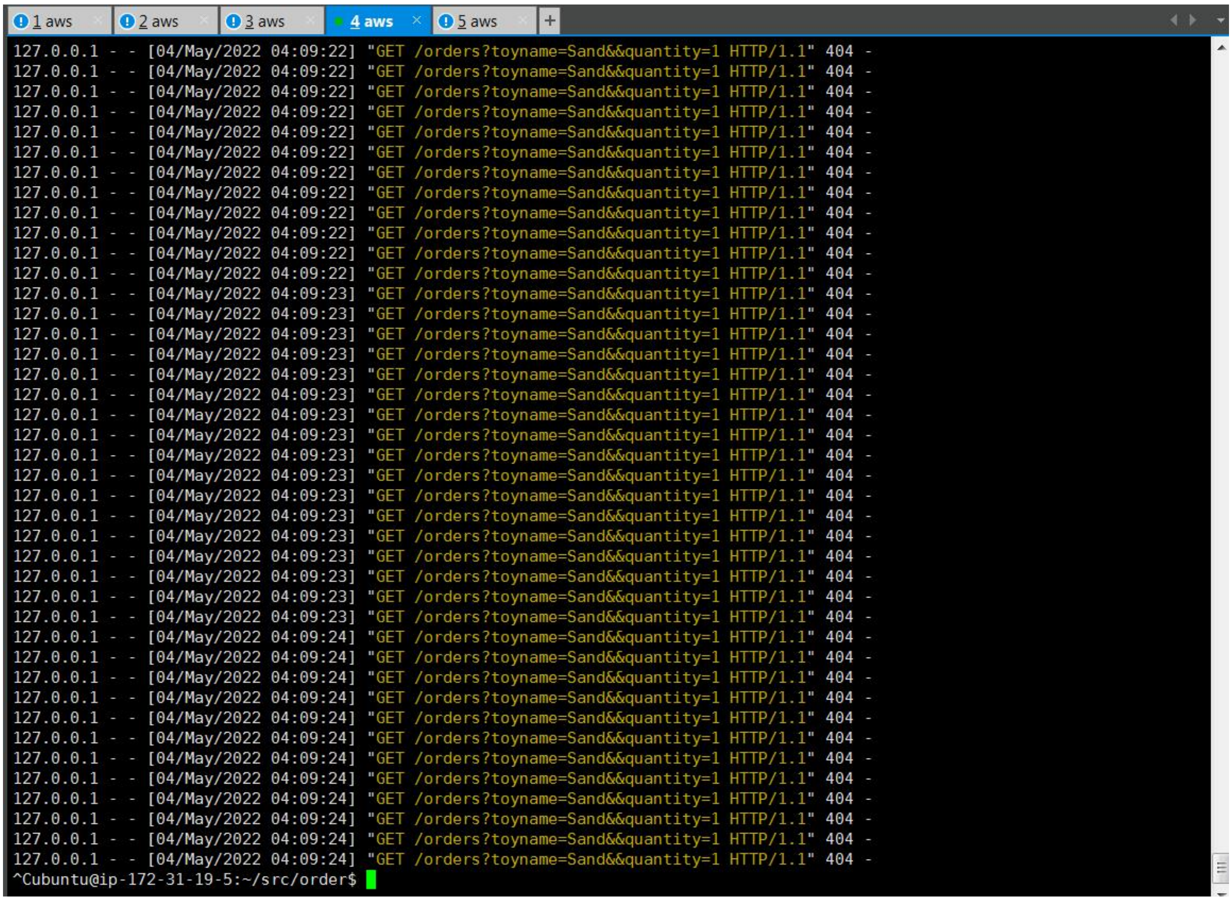

```
1 aws | 2 aws | 3 aws | 4 aws | 5 aws
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:22] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:23] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
127.0.0.1 - - [04/May/2022 04:09:24] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 404 -
^Cubuntu@ip-172-31-19-5:~/src/order$
```

New leader notification: since leader is crashed, front end performed the **leader election**, and notify other nodes who is the new leader.

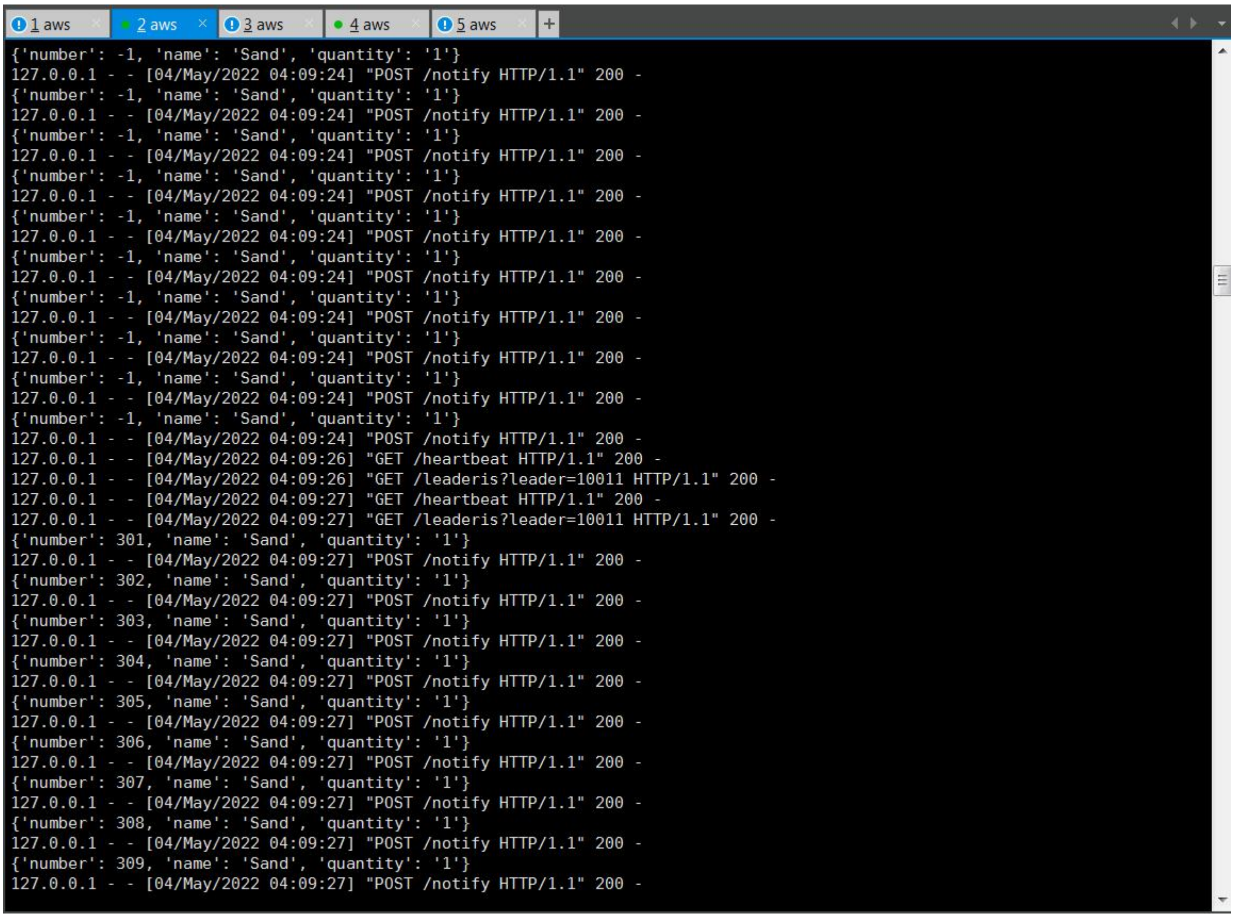

```
1 aws | 2 aws | 3 aws | 4 aws | 5 aws
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:26] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:26] "GET /leaderis?leader=10011 HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:27] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:27] "GET /leaderis?leader=10011 HTTP/1.1" 200 -
{'number': 301, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 302, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 303, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 304, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 305, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 306, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 307, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 308, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
{'number': 309, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:27] "POST /notify HTTP/1.1" 200 -
```

```
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
{'number': -1, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:24] "POST /notify HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:26] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:26] "GET /leaderis?leader=10011 HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:27] "GET /heartbeat HTTP/1.1" 200 -
127.0.0.1 - - [04/May/2022 04:09:27] "GET /leaderis?leader=10011 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
10012 failed!
127.0.0.1 - - [04/May/2022 04:09:27] "GET /orders?toyname=Sand&&quantity=1 HTTP/1.1" 200 -
```

Restart the follower with id = 3: restart the node with **port = 10012** & **id=3**

```
: '1'}, {'number': '430', 'name': 'Sand', 'quantity': '1'}, {'number': '431', 'name': 'Sand', 'quantity': '1'}, {'number': '
432', 'name': 'Sand', 'quantity': '1'}, {'number': '433', 'name': 'Sand', 'quantity': '1'}, {'number': '434', 'name': 'Sand'
, 'quantity': '1'}, {'number': '435', 'name': 'Sand', 'quantity': '1'}, {'number': '436', 'name': 'Sand', 'quantity': '1'},
{'number': '437', 'name': 'Sand', 'quantity': '1'}, {'number': '438', 'name': 'Sand', 'quantity': '1'}, {'number': '439', 'n
ame': 'Sand', 'quantity': '1'}, {'number': '440', 'name': 'Sand', 'quantity': '1'}, {'number': '441', 'name': 'Sand', 'quant
ity': '1'}, {'number': '442', 'name': 'Sand', 'quantity': '1'}, {'number': '443', 'name': 'Sand', 'quantity': '1'}, {'number
': '444', 'name': 'Sand', 'quantity': '1'}, {'number': '445', 'name': 'Sand', 'quantity': '1'}, {'number': '446', 'name': 'S
and', 'quantity': '1'}, {'number': '447', 'name': 'Sand', 'quantity': '1'}, {'number': '448', 'name': 'Sand', 'quantity': '1
'}, {'number': '449', 'name': 'Sand', 'quantity': '1'}]
10012   3
 * Serving Flask app 'order_server' (lazy loading)
 * Environment: production
   WARNING: This is a development server. Do not use it in a production deployment.
   Use a production WSGI server instead.
 * Debug mode: off
 * Running on all addresses.
   WARNING: This is a development server. Do not use it in a production deployment.
 * Running on http://172.31.19.5:10012/ (Press CTRL+C to quit)
{'number': 450, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 451, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 452, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 453, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 454, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 455, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 456, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 457, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 458, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 459, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 460, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
{'number': 461, 'name': 'Sand', 'quantity': '1'}
127.0.0.1 - - [04/May/2022 04:09:38] "POST /notify HTTP/1.1" 200 -
```

Question 2

Can the clients notice the failures (either during order requests or the final order checking phase) or are they transparent to the clients? Do all the order service replicas end up with the same database file?

Solution:

First of all, we are specifying two important factors:

Total latency: If the number of latency is too high, clients will notice the crash.

The number of lines in order log file: In our design, if an order is placed successfully, it will be recorded into order log with an increment unique ID. If a buy is unsuccessful (invalid toy name or out of stock), it will be recorded into order log with ID = -1. Hence, in our test case, if the number of lines in order log is much less than **1000**, it indicates that there are many TCP packets loss during the communication. The clients cannot even receive a successful / unsuccessful response. So we could say the clients can notice the crash in this scenario.

For the test case described previously, the results of those two factors are:

Total latency seen by clients:

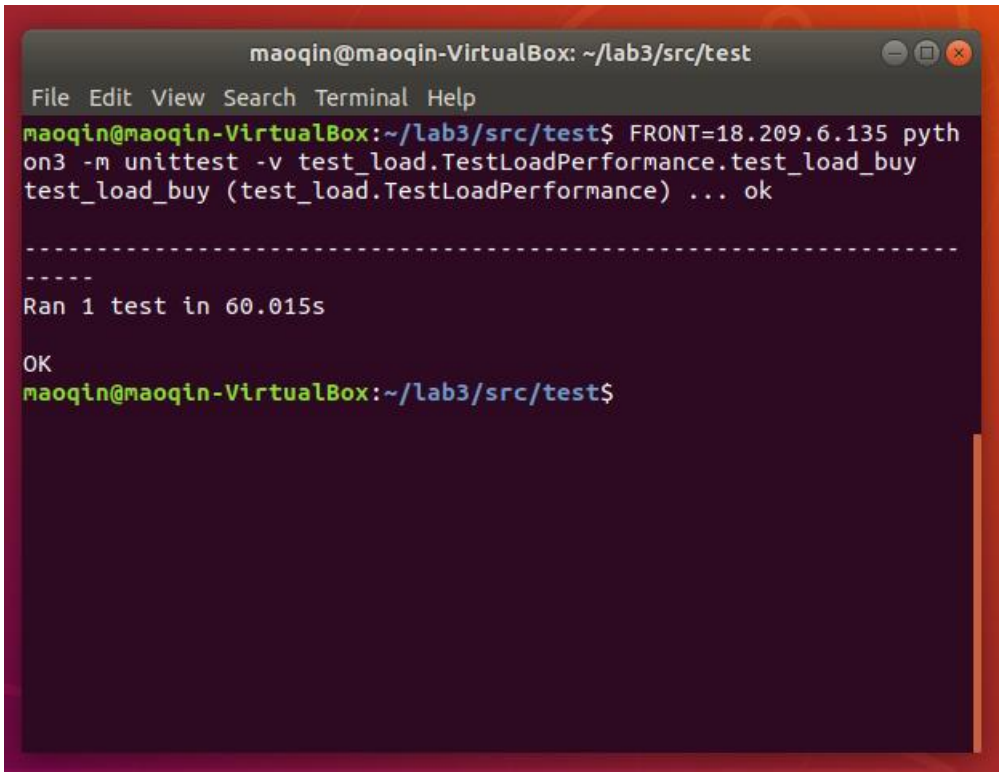

Order log at each order server:

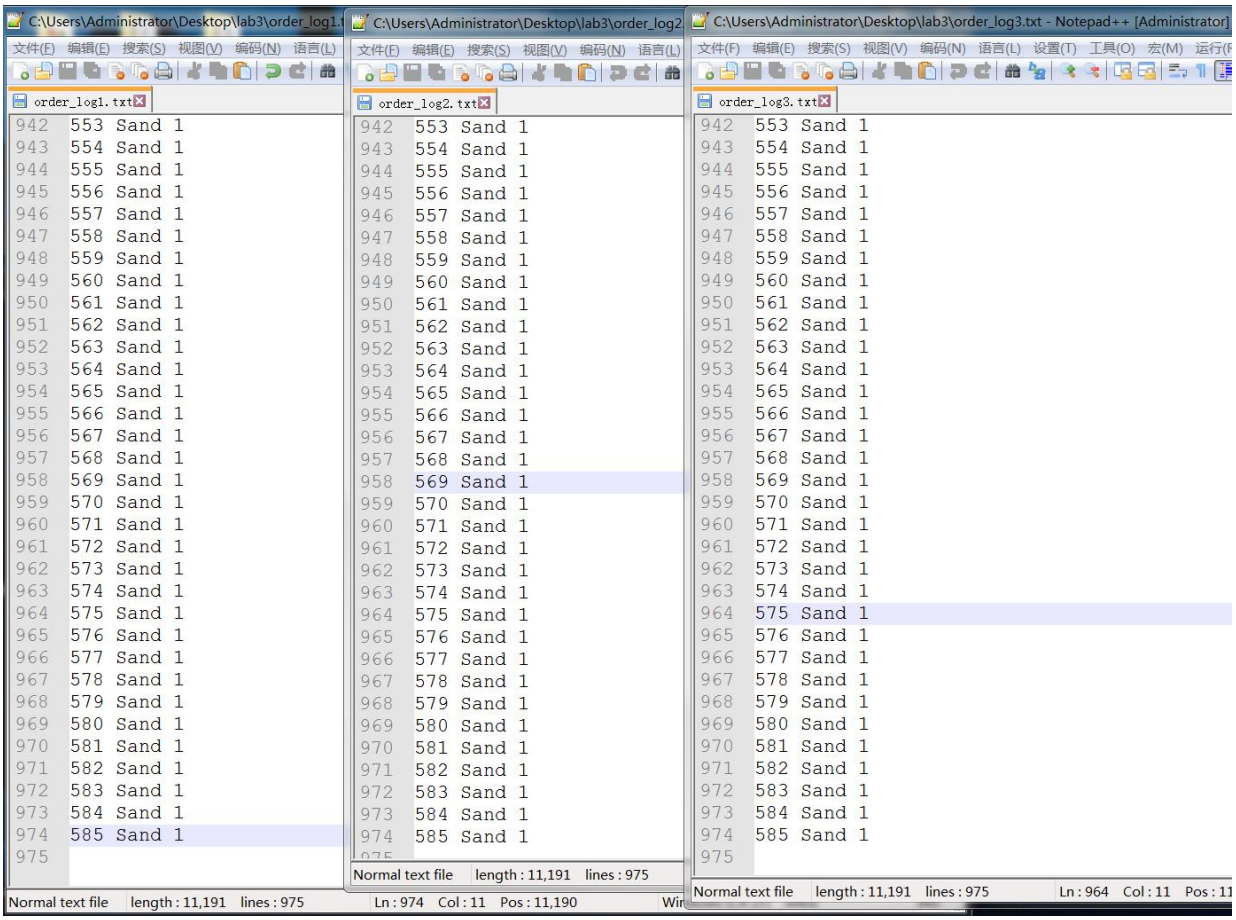

In order to evaluate in what degree the clients can notice the failure, we do the same experiment **without artificial crashes**.

Total latency seen by clients without artificial crashes:

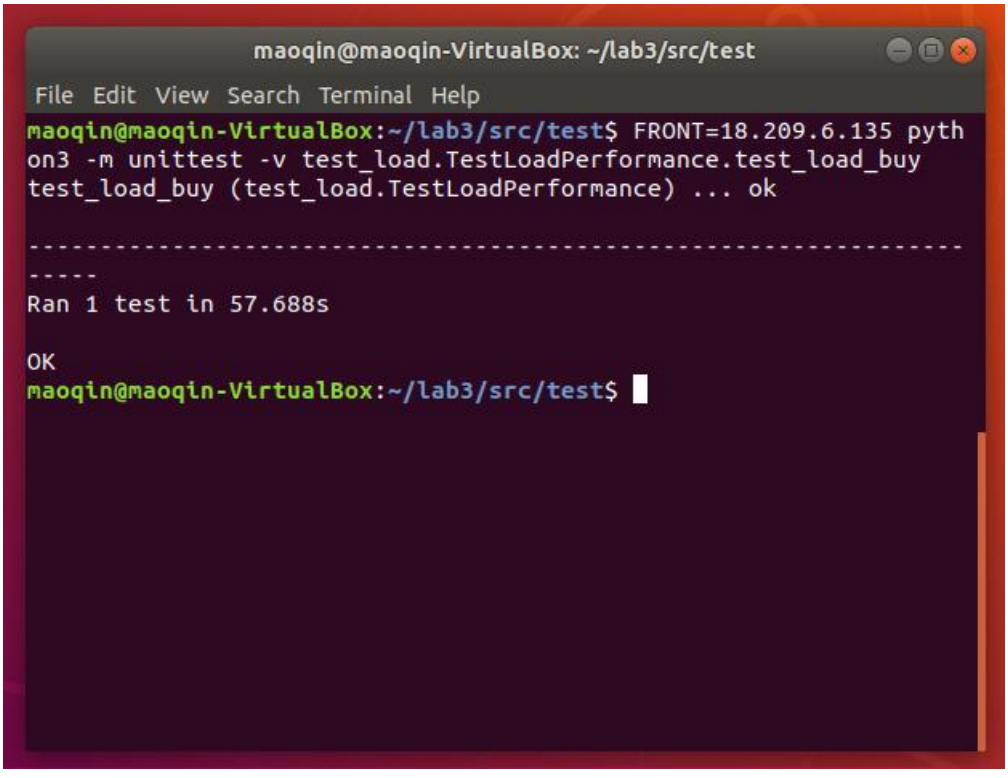

Order log at each order server without artificial crashes:

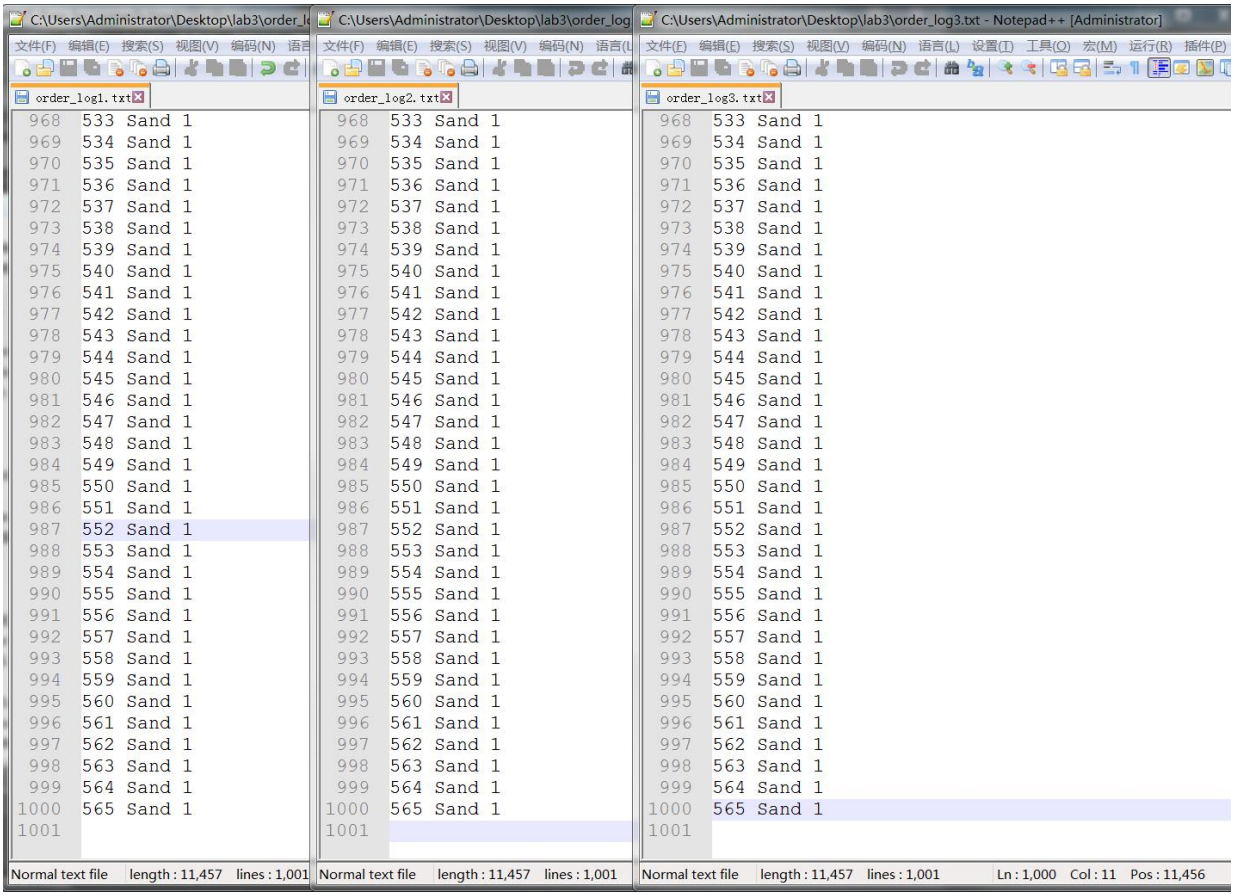

	Total Latency / s	*Packet Loss*
Artificial Crash	*60.015*	*26*
Normal Case	*57.688*	*0*

As you can see, the difference of total latency between crashed case and normal case is really small. And in crashed case, the **packet loss rate** is:

$$packet\ loss\ rate\ =\ 26/1000\ =\ 0.026$$

On balance, we can give our **conclusion** as follows:

1) Clients almost cannot notice the failures (either during order requests or the final order checking phase). The crashes are transparent to the clients

2) As shown above, all the order servers end up with the same database file.